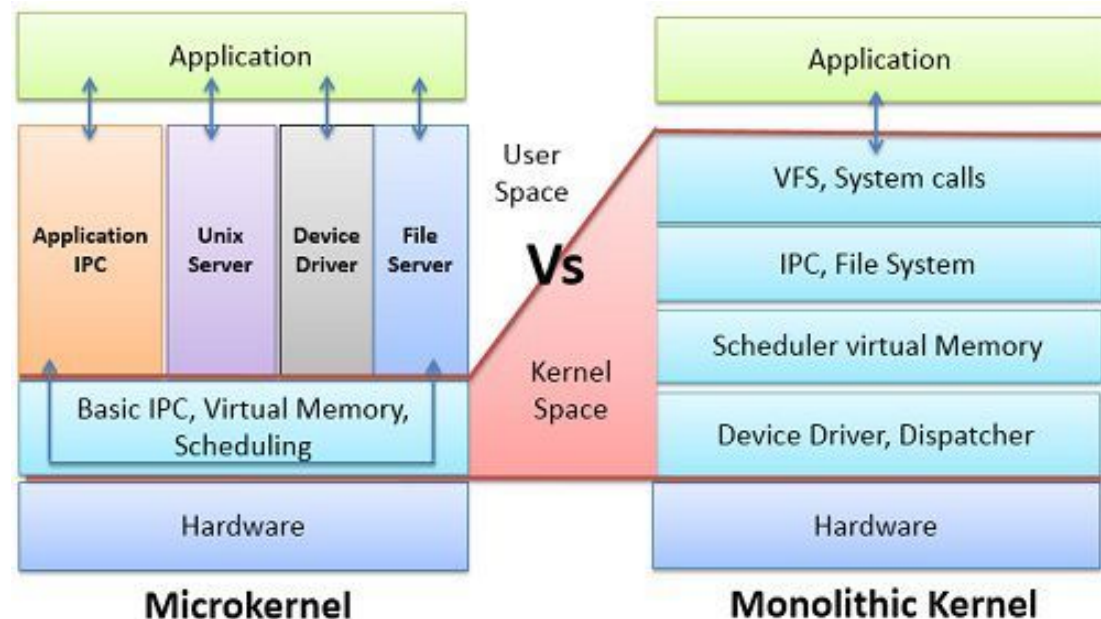


University of Bucharest  
Faculty of Mathematics and Computer Science

# SeL4 verification in Isabelle/HOL

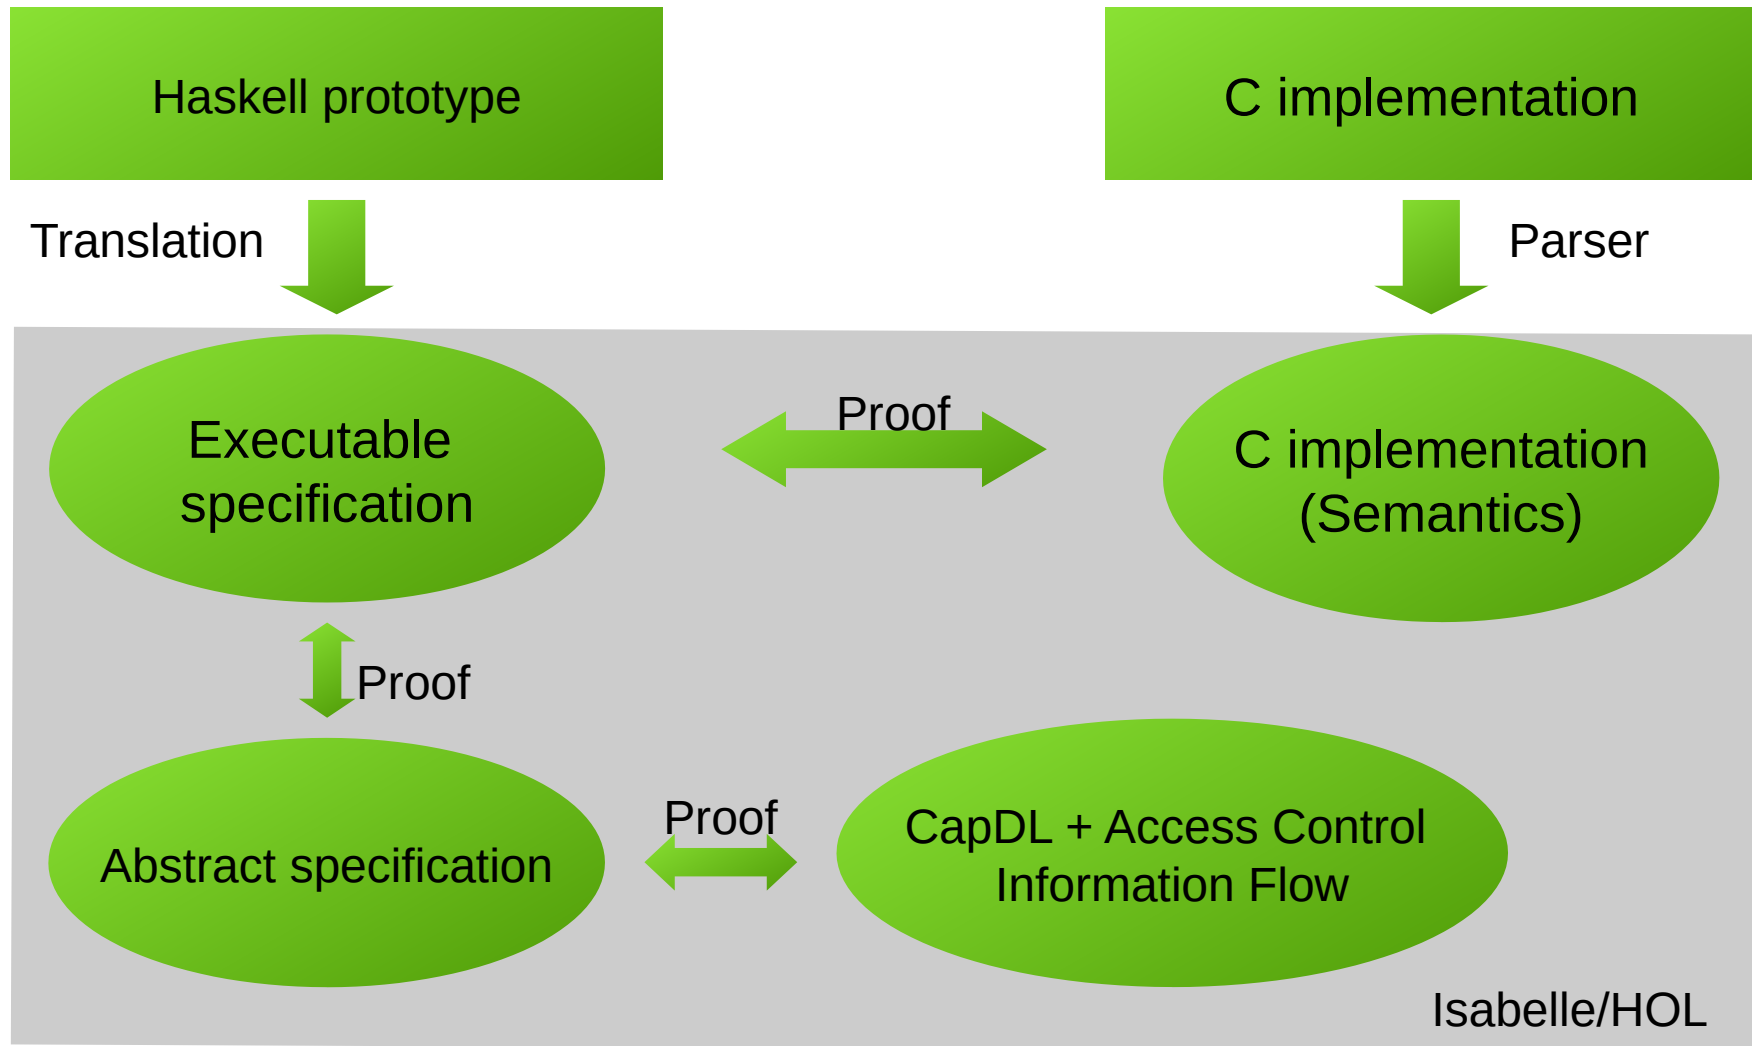
# What is SeL4

- Operating system microkernel (vs monolithic kernels)
- Developed by National Information and Communications Technology Australia – start in 2009
- Short syscalls, interrupts are disabled while in kernel mode
- Capability-based access control
  - Capability = “fat pointer” immutable + access rights
- Functional correctness proof



<https://techdifferences.com/difference-between-microkernel-and-monolithic-kernel.html>

# How kernel verification is done



# OS Verification projects

Project	Highest Level	Lowest Level	Specs	Proofs	Prover	Approach	Year
UCLA Secure Unix	security model	Pascal	90%	20%	XIVUS	Alphard	(?) - 1980
PSOS	application level	source code	17 layers	0%	SPECIAL	HDM	1973 - 1983
KIT	isolated tasks	assembly	100%	100%	Boyer Moore	interpreter equivalence	(?) - 1987
VFiasco/ Rodin	does not crash	C++	70%	0%	PVS	semantic compiler	2001 - 2008
EROS/ Coyotos	security model	BitC	security model	0%	ACL2 (?)	language based	2004 - (?)
Verisoft	application level	gate level	100%	75%	Isabelle	fully pervasive	2004 - (2008)
L4.verified	security model	C/assembly	100%	70%	Isabelle	performance, production code	2005 - (2008)

G. Klein - Operating System Verification - An Overview, NICTA

# Sel4 based OSes



[https://genode.org/documentation/articles/sel4\\_part\\_1](https://genode.org/documentation/articles/sel4_part_1)



+

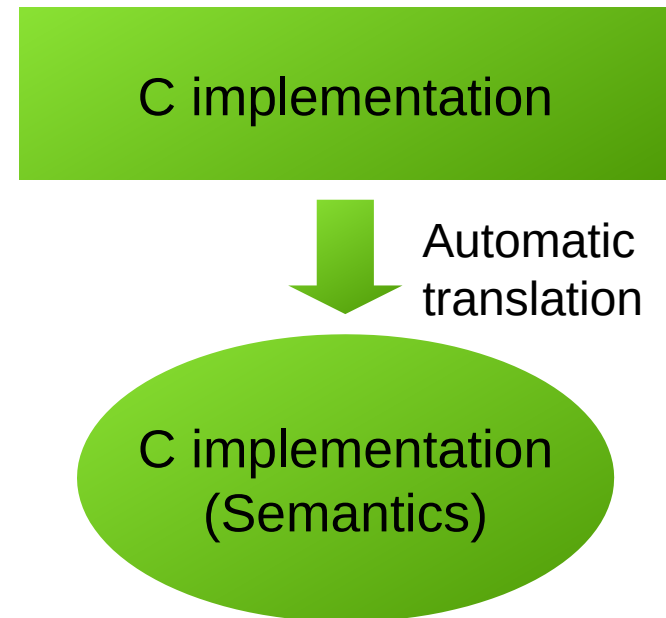


=

**Robigalia OS**  
<https://rbg.systems/>

# C-to-Isabelle Parser

- Implemented by Michael Norrish
- Translates C into its representation in Simpl (imperative programs in Isabelle/HOL)
- Only a subset of C language is used
- Proof<sup>\*</sup> for a large number of programs
- that output matches the binary
- compiled with gcc



<sup>\*</sup>T. Sewell, M. Myreen, G. Klein. *Translation Validation for a Verified OS Kernel*

# C-to-Isabelle Parser

```
int max(int a, int b) {  
  if (a <= b)  
    return b;  
  return a;  
}
```

$\text{max } a \ b \equiv$   
 $\text{if } a \leq b \text{ then } b \text{ else } a$

```
TRY  
  IF { 'a ≤s 'b } THEN  
    'ret__int ::= 'b;;  
    'global_exn_var ::= Return;;  
    THROW  
  ELSE  
    SKIP  
  FI;;  
  'ret__int ::= 'a;;  
  'global_exn_var ::= Return;;  
  THROW;;  
  GUARD DontReach ∅  
  SKIP  
CATCH  
  SKIP  
END
```

D. Greenaway *Automated proof-producing abstraction of C code*

# Deeply and shallowly embedded

$$2 + 2 = 4$$

**Deep embedding** – representation of program *structure*:

*C-to-Isabelle parser*

**Shallow embedding** – representation of program *semantics*:

*AutoCorres tool*

In order to prove two programs are equivalent, we need a *shallow embedded* representation.

AutoCorres generates that model through monadic representation bound to program state.



# C subset and limitations

- Explicit guards against undefined behaviour (division by zero, dereferencing null pointer)
- Allowed only expressions without side effects
- Functions that return values may be called only as the right-hand side of an assignment
- No function pointers, *goto* or *switch* statements, unions or bit-fields also unsupported

```
int i = 0;
int a[2] = {0, 0};
int f(void)
{
    i++;
    return i;
}

int main()
{
    a[i]=f();
    return a[0];
}
```

# Octrng driver demo

```
void
octrng_rnd(void)
{
    unsigned int value;

    rand_value = get_register(OCTRNG_ENTROPY_REG);
    add_task(octrng_rnd, 10);
}
```

A function from driver adaptation

```
octrng_rnd' ≡ do ret' <- get_register' 0;
               modify (rand_value_'_update (λa. ret'));
               add_task' (PTR(unit) (symbol_table 'octrng_rnd')) 10
od
```

The representation of function in Isabelle

# Octrng driver demo

```
octrng_rnd'  $\equiv$  do ret' <- get_register' 0;  
                modify (rand_value_'_update ( $\lambda$ a. ret'));  
                add_task' (PTR(unit) (symbol_table 'octrng_rnd')) 10  
            od
```

The representation of function in Isabelle

```
thm octrng_rnd'_def  
lemma octrng_rnd:  
  "{  $\lambda$ s. timer_' s = a  $\wedge$  running_tasks_' s < MAX_QUEUE  $\wedge$  current_task_' s < MAX_QUEUE  $\wedge$   
    control_addr_C (rng_regs_' s)  $\&\&$  OCTRNG_ENABLE_OUTPUT  $\neq$  0  $\wedge$   
    control_addr_C (rng_regs_' s)  $\&\&$  OCTRNG_ENABLE_ENTROPY  $\neq$  0 }  
  octrng_rnd'  
  {  $\lambda$ s. rand_value_' s = a }!"  
unfolding octrng_rnd'_def  
unfolding get_register'_def add_task'_def  
unfolding get_time'_def  
unfolding MAX_QUEUE_def  
unfolding OCTRNG_ENABLE_OUTPUT_def OCTRNG_ENABLE_ENTROPY_def  
apply (wp; auto)+  
done
```

Proof of lemma stating that  
octrng\_rnd function returns current time

# Bibliography

- D. Greenaway, “*Automated proof-producing abstraction of C code*”, Ph.D. Dissertation, School of Computer Science and Engineering, University of NSW, Sydney, Australia, (2014);
- G. Heiser, “*The seL4 Microkernel – An Introduction*”, White paper, The seL4 Foundation, rev 1.2 from 2020-06-10;
- H. Tuch, G. Klein, M. Norrish, “*Types, Bytes, and Separation Logic*” in Proceedings of the 34th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, Nice, France, 97–108, (2007);
- D. Greenaway, J. Andronick, K. Gerwin, “*Bridging the Gap: Automatic Verified Abstraction of C*” (2012);
- G. Klein, J. Andronick, K. Elphinstone, T. Murray, T. Sewel, R. Kolanski, G. Heiser, NICTA and UNSW, “*Comprehensive Formal Verification of an OS Microkernel*”, Sydney, Australia, ACM Transactions on Computer Systems, Vol. 32, No. 1, Art. 2, (2014).