

C08 – SMT Solvers, Symbolic execution

Program Verification

FMI · Denisa Diaconescu · Spring 2022

How SMT solvers works?

Symbolic execution

How SMT solvers works?

The SMT problem

The SMT problem:

Given a first-order logic formula, with symbols from (possibly several) theories, does it have a model?

Is the formula satisfiable? If so, how?

The SAT problem is a special case, in which

- the formula is quantifier-free, without function symbols or equality
- no theories are used

SAT solving algorithms are an important ingredient in SMT solvers

First-order theories

- Whereas the language of SAT solvers is Boolean logic, the language of SMT solvers is **first-order logic**.
- **First-order theories** allow us to capture structures which are used by programs (e.g., arrays, integers) and enable reasoning about them.
- **Validity in first order logic (FOL) is undecidable!**
 - **Lambda calculus** – Alonzo Church (1936)
 - **Turing machines** – Alan Turing (1937)
 - **Recursive functions** – Kurt Gödel (1934) and Stephen Kleene (1936)
- Validity in particular first order theories is (sometimes) decidable.

Combine **propositional satisfiability** search techniques with solvers for **specific first-order theories**:

- Linear arithmetic
- Bit vectors
- Arrays
- ...



There are two main approaches for SMT solvers:

- The eager approach
 - Tries to find ways of encoding an entire SMT problem into SAT.
 - There are a variety of techniques
 - For some theories, this works quite well.

There are two main approaches for SMT solvers:

- The eager approach
 - Tries to find ways of encoding an entire SMT problem into SAT.
 - There are a variety of techniques
 - For some theories, this works quite well.
- The lazy approach
 - Tries to combine SAT and theory reasoning.
 - The basis for most modern SMT solvers.

SMT: The Big Questions

1. How to solve conjunctions of literals in a theory?
 - Use a Theory solver
2. How to combine a theory solver and a SAT solver to reason about arbitrary formulas?
 - The DPLL(T) framework
3. How to combine theory solvers for several theories?
 - The Nelson-Oppen method and its variants

- Given a theory T , a Theory solver for T takes as input a set (interpreted as an implicit conjunction) φ of literals and determines whether φ is T -satisfiable.
 - φ is T -satisfiable if there is some model \mathcal{M} of T such that φ holds in \mathcal{M} .
- In order to integrate a Theory solver into a modern SMT solver, it is helpful if the Theory solver can do more than just check satisfiability.

Characteristics of Theory solvers

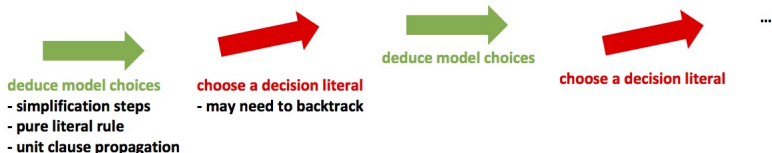
Some desirable characteristics of Theory solvers include:

- **Incrementality** – easy to add new literals or backtrack to a previous state
- **Layered/Lazy** – able to detect simple inconsistencies quickly, able to detect difficult inconsistencies eventually
- **Equality Propagating** – if Theory solvers can detect when two terms are equivalent, this greatly simplifies theory combination
- **Model Generating** – when reporting T -satisfiable, the Theory solver also provides a concrete value for each variable or function symbol
- **Proof Generating** – when reporting T -unsatisfiable, the Theory solver also provides a checkable proof

Propositional Abstraction

- An **atom** is a formula without propositional connectives or quantifiers
 - depending on the signature $f(a) = b, m * n \leq 42$ could be atoms; 42 is not
 - a propositional atom is an uninterpreted constant symbol of sort Bool
- A (first-order) **literal** is an atom or its negation
- For a given signature Σ , we define a signature Σ^P containing only:
 - the **propositional Σ -atoms**
 - a **fresh propositional atom** for each non-propositional Σ -atom
- We then fix an injective mapping from the non-propositional Σ -atoms to the Σ^P -atoms.
- For a Σ -formula φ , the formula φ^P is the **propositional abstraction** of φ , given by replacing all non-propositional Σ -atoms in φ with their image under this mapping.
- An Σ -formula φ is **propositional unsatisfiable** if $\varphi^P \models \perp$.
- An Σ -formula φ **propositionally entails** an Σ -formula ψ if $\varphi^P \models \psi^P$.
 - Note that $\varphi^P \models \psi^P$ implies $\varphi \models \psi$, but **not necessarily vice-versa**.

Recall DPLL/CDCL Algorithms

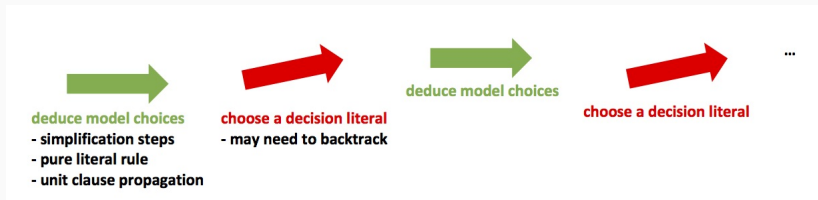


...until...

- conflict reached
 - backtrack - try flipping a decision literal
 - (if CDCL) learn new clause, back-jump
- model found
 - return the model

Adapting DPLL to DPLL(T)

Run DPLL on the **propositional abstraction** φ^P of the T -input formula φ



...until...

- **conflict reached:** backtrack/jump, learn clauses as usual
- **model found** (represented by a set Γ of literals)
 - It is not necessarily a T -model!
 - Ask theory solver: is Γ T -satisfiable?
 - If yes, we are done.
 - If no, backtrack in the original search.
 - (CDCL) get a T -unsatisfiable subset for clause learning/back-jumping

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4]$ (i.e. $[[1], [-2, 3], [-4]]$)

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4]$ (i.e. $[[1], [-2, 3], [-4]]$)
- SAT solver returns model $[1, \neg 2, \neg 4]$

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4]$ (i.e. $[[1], [-2, 3], [-4]]$)
- SAT solver returns model $[1, \neg 2, \neg 4]$
- Theory solver detects $[1, \neg 2]$ **T-unsat**

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4]$ (i.e. $[[1], [-2, 3], [-4]]$)
- SAT solver returns model $[1, \neg 2, \neg 4]$
- Theory solver detects $[1, \neg 2]$ **T-unsat**
- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4, \neg 1 \vee 2]$

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4]$ (i.e. $[[1], [-2, 3], [-4]]$)
- SAT solver returns model $[1, \neg 2, \neg 4]$
- Theory solver detects $[1, \neg 2]$ **T-unsat**
- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4, \neg 1 \vee 2]$
- SAT solver returns model $[1, 2, 3, \neg 4]$

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4]$ (i.e. $[[1], [-2, 3], [-4]]$)
- SAT solver returns model $[1, \neg 2, \neg 4]$
- Theory solver detects $[1, \neg 2]$ **T-unsat**
- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4, \neg 1 \vee 2]$
- SAT solver returns model $[1, 2, 3, \neg 4]$
- Theory solver detects $[1, 3, \neg 4]$ **T-unsat**

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4]$ (i.e. $[[1], [-2, 3], [-4]]$)
- SAT solver returns model $[1, \neg 2, \neg 4]$
- Theory solver detects $[1, \neg 2]$ **T-unsat**
- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4, \neg 1 \vee 2]$
- SAT solver returns model $[1, 2, 3, \neg 4]$
- Theory solver detects $[1, 3, \neg 4]$ **T-unsat**
- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4, \neg 1 \vee 2, \neg 1 \vee \neg 3 \vee 4]$

Adapting DPLL to DPLL(T)

Example

$$\underbrace{g(a) = c}_1 \wedge \underbrace{(f(g(a)) \neq f(c))}_{\neg 2} \vee \underbrace{g(a) = d}_3 \wedge \underbrace{c \neq d}_{\neg 4}$$

- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4]$ (i.e. $[[1], [-2, 3], [-4]]$)
- SAT solver returns model $[1, \neg 2, \neg 4]$
- Theory solver detects $[1, \neg 2]$ **T-unsat**
- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4, \neg 1 \vee 2]$
- SAT solver returns model $[1, 2, 3, \neg 4]$
- Theory solver detects $[1, 3, \neg 4]$ **T-unsat**
- Call SAT solver with input $[1, \neg 2 \vee 3, \neg 4, \neg 1 \vee 2, \neg 1 \vee \neg 3 \vee 4]$
- SAT solver detects **unsat**

Theory combination

- Given a theory T , a **Theory solver** for T takes as input a set (interpreted as an implicit conjunction) φ of literals and determines whether φ is T -satisfiable.
- We are often interested in using **two or more theories at the same time**.
- **Can we combine two theory solvers to get a theory solver for the combined theory?**

Theory combination

- Given a theory T , a **Theory solver** for T takes as input a set (interpreted as an implicit conjunction) φ of literals and determines whether φ is T -satisfiable.
- We are often interested in using **two or more theories at the same time**.
- **Can we combine two theory solvers to get a theory solver for the combined theory?**

Example

The following formula uses both T_E and T_Z

$$\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$$

The Nelson-Oppen Method

A very general method for combining theory solvers is the Nelson-Oppen method.

This method is applicable when:

1. The signatures Σ_i are disjoint.
2. The theories T_i are stably-infinite.
 - A Σ -theory T is stably-infinite if every T -satisfiable quantifier-free Σ -formula is satisfiable in an infinite model.
3. The formulas to be tested for satisfiability are conjunctions of quantifier-free literals.

Extensions exist that can relax each of these restrictions in some cases.

The Nelson-Oppen Method

Some definitions:

- A member of Σ_i is an *i*-symbol.
- A term t is an *i*-term if it starts with an *i*-symbol.
- An atomic *i*-formula is
 - an application of an *i*-predicate,
 - an equation whose lhs is an *i*-term, or
 - an equation whose lhs is a variable and whose rhs is an *i*-term
- An *i*-literal is an atomic *i*-formula or the negation of one.
- An occurrence of a term t in either an *i*-term or an *i*-literal is *i*-alien if it is a j -term with $i \neq j$ and all of its super-terms (if any) are *i*-terms.
- An expression is pure if it contains only variables and *i*-symbols for some i .

Conversion to separate form

Given a conjunction of literals φ , we want to convert it into a **separate form**: a T -equisatisfiable conjunction of literals $\varphi_1 \wedge \varphi_2 \wedge \dots \wedge \varphi_n$ where each φ_i is a Σ_i -formula.

We have the following **algorithm**:

1. Let ψ be **some literal** in φ .
2. If ψ is a **pure i -literal**, for some i , remove ψ from φ and add ψ to φ_i .
If φ is **empty then stop**; otherwise **goto step 1**.
3. Otherwise, ψ is an **i -literal** for some i .
Let t be a **term occurring i -alien** in ψ .
Replace t in φ with a **new variable** z and add $z = t$ to φ .
Goto step 1.

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$$

We convert φ to a separate form:

- $\varphi_E := ?$
- $\varphi_Z := ?$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$$

We convert φ to a separate form:

- $\varphi_E := ?$
- $\varphi_Z := 1 \leq x$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$$

We convert φ to a separate form:

- $\varphi_E := ?$
- $\varphi_Z := 1 \leq x \wedge x \leq 2$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$$

We convert φ to a separate form:

- $\varphi_E := ?$
- $\varphi_Z := 1 \leq x \wedge x \leq 2$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(y) \wedge f(x) \neq f(2) \wedge y = 1$$

We convert φ to a separate form:

- $\varphi_E := ?$
- $\varphi_Z := 1 \leq x \wedge x \leq 2$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = \cancel{1 \leq x \wedge x \leq 2} \wedge f(x) \neq f(y) \wedge f(x) \neq f(2) \wedge y = 1$$

We convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(y) \wedge f(x) \neq f(2) \wedge y = 1$$

We convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(y) \wedge f(x) \neq f(z) \wedge y = 1 \wedge z = 2$$

We convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(y) \wedge f(x) \neq f(z) \wedge y = 1 \wedge z = 2$$

We convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2$

Conversion to separate form

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(y) \wedge f(x) \neq f(z) \wedge y = 1 \wedge z = 2$$

We convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1$

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(y) \wedge f(x) \neq f(z) \wedge y = 1 \wedge z = 2$$

We convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi = 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(y) \wedge f(x) \neq f(z) \wedge y = 1 \wedge z = 2$$

We convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

The Nelson-Oppen Method

- As each φ_i is a Σ_i -formula, we can run a Theory solver Sat_i for each φ_i .
- If any Sat_i reports that φ_i is unsatisfiable, then φ is unsatisfiable.
- The converse is not true in general!
- We need a way for the decision procedures to communicate with each other about shared variables.
- If S is a set of terms and \sim is an equivalence relation on S , then the arrangement of S induced by \sim is

$$Ar_{\sim} = \{x = y \mid x \sim y\} \cup \{x \neq y \mid x \not\sim y\}$$

The Nelson-Oppen Method

Suppose that T_1 and T_2 are theories with disjoint signatures Σ_1 and Σ_2 .

Let $T = \bigcup T_i$ and $\Sigma = \bigcup \Sigma_i$.

Given a Σ -formula φ and decision procedures Sat_1 and Sat_2 for T_1 and T_2 , we wish to determine if φ is T -satisfiable.

The non-deterministic Nelson-Oppen algorithm:

1. Convert φ to its **separate form** $\varphi_1 \wedge \varphi_2$.
2. Let S be the **set of variables shared** between φ_1 and φ_2 .
Guess an **equivalence relation** \sim on S .
3. **Run** Sat_1 on $\varphi_1 \cup Ar_\sim$.
4. **Run** Sat_2 on $\varphi_2 \cup Ar_\sim$.

The Nelson-Oppen Method

Suppose that T_1 and T_2 are theories with disjoint signatures Σ_1 and Σ_2 .

Let $T = \bigcup T_i$ and $\Sigma = \bigcup \Sigma_i$.

Given a Σ -formula φ and decision procedures Sat_1 and Sat_2 for T_1 and T_2 , we wish to determine if φ is T -satisfiable.

The non-deterministic Nelson-Oppen algorithm:

1. Convert φ to its **separate form** $\varphi_1 \wedge \varphi_2$.
2. Let S be the **set of variables shared** between φ_1 and φ_2 .
Guess an **equivalence relation** \sim on S .
3. **Run** Sat_1 on $\varphi_1 \cup Ar_\sim$.
4. **Run** Sat_2 on $\varphi_2 \cup Ar_\sim$.

If there exists an equivalence relation \sim such that both Sat_1 and Sat_2 succeed, then φ is **T -satisfiable**.

If no such equivalence relation exists, then φ is **T -unsatisfiable**.

The Nelson-Oppen Method

Suppose that T_1 and T_2 are theories with disjoint signatures Σ_1 and Σ_2 .

Let $T = \bigcup T_i$ and $\Sigma = \bigcup \Sigma_i$.

Given a Σ -formula φ and decision procedures Sat_1 and Sat_2 for T_1 and T_2 , we wish to determine if φ is T -satisfiable.

The non-deterministic Nelson-Oppen algorithm:

1. Convert φ to its **separate form** $\varphi_1 \wedge \varphi_2$.
2. Let S be the **set of variables shared** between φ_1 and φ_2 .
Guess an **equivalence relation** \sim on S .
3. **Run** Sat_1 on $\varphi_1 \cup Ar_\sim$.
4. **Run** Sat_2 on $\varphi_2 \cup Ar_\sim$.

If there exists an equivalence relation \sim such that both Sat_1 and Sat_2 succeed, then φ is **T -satisfiable**.

If no such equivalence relation exists, then φ is **T -unsatisfiable**.

The generalization to more than two theories is straightforward.

The Nelson-Oppen Method

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$$

We first convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

The Nelson-Oppen Method

Example

Consider the following $\Sigma_E \cup \Sigma_Z$:

$$\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$$

We first convert φ to a separate form:

- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

The shared variables are $\{x, y, z\}$.

There are 5 possible arrangements based on equivalence classes of x, y , and z (see *Bell number*).

Example

- $\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$
- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

1. $\{x = y, x = z, y = z\}$
2. $\{x = y, x \neq z, y \neq z\}$
3. $\{x \neq y, x = z, y \neq z\}$
4. $\{x \neq y, x \neq z, y = z\}$
5. $\{x \neq y, x \neq z, y \neq z\}$

The Nelson-Oppen Method

Example

- $\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$
- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

1. $\{x = y, x = z, y = z\}$ inconsistent with T_E
2. $\{x = y, x \neq z, y \neq z\}$
3. $\{x \neq y, x = z, y \neq z\}$
4. $\{x \neq y, x \neq z, y = z\}$
5. $\{x \neq y, x \neq z, y \neq z\}$

The Nelson-Oppen Method

Example

- $\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$
- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

1. $\{x = y, x = z, y = z\}$ inconsistent with T_E
2. $\{x = y, x \neq z, y \neq z\}$ inconsistent with T_E
3. $\{x \neq y, x = z, y \neq z\}$
4. $\{x \neq y, x \neq z, y = z\}$
5. $\{x \neq y, x \neq z, y \neq z\}$

The Nelson-Oppen Method

Example

- $\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$
- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

1. $\{x = y, x = z, y = z\}$ inconsistent with T_E
2. $\{x = y, x \neq z, y \neq z\}$ inconsistent with T_E
3. $\{x \neq y, x = z, y \neq z\}$ inconsistent with T_E
4. $\{x \neq y, x \neq z, y = z\}$
5. $\{x \neq y, x \neq z, y \neq z\}$

The Nelson-Oppen Method

Example

- $\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$
- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

- | | |
|---------------------------------------|-------------------------|
| 1. $\{x = y, x = z, y = z\}$ | inconsistent with T_E |
| 2. $\{x = y, x \neq z, y \neq z\}$ | inconsistent with T_E |
| 3. $\{x \neq y, x = z, y \neq z\}$ | inconsistent with T_E |
| 4. $\{x \neq y, x \neq z, y = z\}$ | inconsistent with T_Z |
| 5. $\{x \neq y, x \neq z, y \neq z\}$ | |

The Nelson-Oppen Method

Example

- $\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$
- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

- | | |
|---------------------------------------|-------------------------|
| 1. $\{x = y, x = z, y = z\}$ | inconsistent with T_E |
| 2. $\{x = y, x \neq z, y \neq z\}$ | inconsistent with T_E |
| 3. $\{x \neq y, x = z, y \neq z\}$ | inconsistent with T_E |
| 4. $\{x \neq y, x \neq z, y = z\}$ | inconsistent with T_Z |
| 5. $\{x \neq y, x \neq z, y \neq z\}$ | inconsistent with T_Z |

The Nelson-Oppen Method

Example

- $\varphi := 1 \leq x \wedge x \leq 2 \wedge f(x) \neq f(1) \wedge f(x) \neq f(2)$
- $\varphi_E := f(x) \neq f(y) \wedge f(x) \neq f(z)$
- $\varphi_Z := 1 \leq x \wedge x \leq 2 \wedge y = 1 \wedge z = 2$

- | | |
|---------------------------------------|-------------------------|
| 1. $\{x = y, x = z, y = z\}$ | inconsistent with T_E |
| 2. $\{x = y, x \neq z, y \neq z\}$ | inconsistent with T_E |
| 3. $\{x \neq y, x = z, y \neq z\}$ | inconsistent with T_E |
| 4. $\{x \neq y, x \neq z, y = z\}$ | inconsistent with T_Z |
| 5. $\{x \neq y, x \neq z, y \neq z\}$ | inconsistent with T_Z |

Conclusion: φ is $T_E \cup T_Z$ -unsatisfiable!

Recall the ingredients:

- Theory solvers for different theories
- Combine a Theory solver and a SAT solver
- Combine Theory solvers for different theories

Symbolic execution

- Symbolic execution is widely used in practice.
- Tools based on symbolic execution have found serious errors and security vulnerabilities in various systems:
 - Networks servers
 - File systems
 - Device drivers
 - Unix utilities
 - Computer vision code
 - ...

- Stanford's KLEE
<http://klee.llvm.org/>
- Nasa's Java PathFinder
<http://javapathfinder.sourceforge.net/>
- Microsoft Research's SAFE
- UC Berkeley's CUTE

At any point during program execution, symbolic execution keeps two formulas:

- symbolic store and
- path constraint

Therefore, at any point in time the symbolic state is described as the conjunction of these two formulas.

The **value of variables** at any moment in time are given by a function

$$\sigma_s : Var \rightarrow Sym$$

- Var is the set of variables
- Sym is a set of **symbolic values**
- σ_s is called a **symbolic store**

Example

$$\sigma_s : x \mapsto x0, y \mapsto y0$$

Arithmetic expression evaluation simply manipulates the symbolic values.

Example

Suppose the symbolic store is $\sigma_s : x \mapsto x_0, y \mapsto y_0$.

Then $z = x + y$ will produce the new symbolic store

$$\sigma'_s : x \mapsto x_0, y \mapsto y_0, z \mapsto x_0 + y_0$$

We literally keep the **symbolic expression** $x_0 + y_0$.

- The analysis keeps a **path constraint** (*pct*) which records the history of all branches taken so far.
- The path constraint is simply a **formula**.
- The formula is typically in a decidable logical fragment without quantifiers.
- At the start of the analysis, the path constraint is **true**.
- Evaluation of **conditionals** affects the path constraint, but not the symbolic store.

Path constraint

Example

Suppose the symbolic store is $\sigma_s : x \mapsto x0, y \mapsto y0$.

Suppose the path constraint is $pct = x0 > 10$.

Let us evaluate `if(x > y + 1) {5: ...}`

At label 5, we will get the symbolic store σ_s . It does not change!

But, at label 5, we will get an **updated path constraint**:

$$pct = x0 > 10 \wedge x0 > y0 + 1$$

Symbolic execution - example

```
int twice(int v) {  
    return 2 * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x,y);  
}
```

Can you find the inputs that make the program reach the ERROR?

Lets execute this example with classic symbolic execution

source: Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.

Symbolic execution - example

```
int twice(int v) {  
    return 2 * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x, y);  
}
```

The read() functions read a value from the input and because we don't know what those read values are, we set the values of x and y to fresh symbolic values called x_0 and y_0

pct is true because so far we have not executed any conditionals

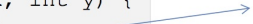
$\sigma_s : \quad x \mapsto x_0,$
 $\quad \quad y \mapsto y_0$

pct : true

source: Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.

Symbolic execution - example

```
int twice(int v) {  
    return 2 * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x, y);  
}
```



$\sigma_s : \begin{array}{l} x \mapsto x0, \\ y \mapsto y0 \\ z \mapsto 2*y0 \end{array}$

pct : true

Here, we simply executed the function twice() and added the new symbolic value for z.

source: Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.

Symbolic execution - example

```
int twice(int v) {  
    return 2 * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x, y);  
}
```

We forked the analysis into 2 paths: the true and the false path. So we **duplicate** the state of the analysis.

This is the result if $x = z$:

$\sigma_s : x \mapsto x_0,$
 $y \mapsto y_0$
 $z \mapsto 2 * y_0$

$pct : x_0 = 2 * y_0$

This is the result if $x \neq z$:

$\sigma_s : x \mapsto x_0,$
 $y \mapsto y_0$
 $z \mapsto 2 * y_0$

$pct : x_0 \neq 2 * y_0$

source: Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.

Symbolic execution - example

```
int twice(int v) {  
    return 2 * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x, y);  
}
```

We can avoid further exploring a path if we know the constraint `pct` is **unsatisfiable**. In this example, both `pct`'s are **satisfiable** so we need to keep exploring both paths.

This is the result if $x = z$:

$\sigma_s : x \mapsto x_0,$
 $y \mapsto y_0$
 $z \mapsto 2 * y_0$

$pct : x_0 = 2 * y_0$

This is the result if $x \neq z$:

$\sigma_s : x \mapsto x_0,$
 $y \mapsto y_0$
 $z \mapsto 2 * y_0$

$pct : x_0 \neq 2 * y_0$

source: Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.

Symbolic execution - example

```
int twice(int v) {  
    return 2 * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x, y);  
}
```

Lets explore the path when $x == z$ is true.
Once again we get 2 more paths.

This is the result if $x > y + 10$:

$\sigma_s : x \mapsto x0,$
 $y \mapsto y0$
 $z \mapsto 2*y0$

$pct : x0 = 2*y0$
 \wedge
 $x0 > y0+10$

This is the result if $x \leq y + 10$:

$\sigma_s : x \mapsto x0,$
 $y \mapsto y0$
 $z \mapsto 2*y0$

$pct : x0 = 2*y0$
 \wedge
 $x0 \leq y0+10$

source: Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.

Symbolic execution - example

```
int twice(int v) {  
    return 2 * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x, y);  
}
```

So the following path reaches **"ERROR"**.

This is the result if $x > y + 10$:

$\sigma_s : x \mapsto x_0,$
 $y \mapsto y_0$
 $z \mapsto 2 * y_0$

$pct : x_0 = 2 * y_0$
 \wedge
 $x_0 > y_0 + 10$

We can now ask the SMT solver for a satisfying assignment to the pct formula.

For instance, $x_0 = 40, y_0 = 20$ is a satisfying assignment. That is, running the program with those concrete inputs triggers the error.

Handling Loops - a limitation

```
int F(unsigned int k) {  
    int sum = 0;  
    int i = 0;  
    for ( ; i < k; i++)  
        sum += i;  
    return sum;  
}
```

- A serious limitation of symbolic execution is **handling unbounded loops**.
- Symbolic execution runs the program for a finite number of paths.
- But what happens if we do not know the bound on a loop?
- **The symbolic execution will keep running forever!**

Handling Loops - bound loops


```
int F(unsigned int k) {  
    int sum = 0;  
    int i = 0;  
    for ( ; i < 2; i++)  
        sum += i;  
    return sum;  
}
```

- A common solution in practice is to **provide some loop bound**.
 - In the above example, we can bound `k` to say 2.
 - This is an example of an **under-approximation**
- Practical symbolic analyzers usually under-approximate as most programs have unknown loop bounds.

Handling Loops - loop invariants

```
int F(unsigned int k) {  
    int sum = 0;  
    int i = 0;  
    for ( ; i < k; i++)  
        sum += i;  
    return sum;  
}
```

loop invariant



- Another solution is to **provide a loop invariant**.
- This technique is rarely used for large programs because it is **difficult to provide** such invariants manually.
- It can also lead to **over-approximation**.

Constraint solving - challenges

- Constraint solving is fundamental to symbolic execution.
- An SMT solver is continuously invoked during analysis.
- Often, the main **roadblock to performance** of symbolic execution engines is the time spent in constraint solving.
- Important features:
 - The SMT solver supports as **many decidable logical fragments** as possible.
 - Some tools use more than one SMT solver.
 - The SMT solver can **solve large formulas quickly**.
 - The symbolic execution engines tries to reduce the burden in calling the SMT solver by **exploring domain specific insights**.

Key optimization - caching

- The analyzer will invoke the SMT solver with **similar formulas**.
- The symbolic execution engine can keep a **map (cache)** of formulas to a satisfying assignment for the formulas.
- When the engine builds a new formula and would like to find a satisfying assignment for that formula, it can **first access the cache**, before calling the SMT solver.

Key optimization - caching

Example

Suppose the cache contains the mapping:

Formula		Solution
$(x + y < 10) \wedge (x > 5)$	\rightarrow	$\{x = 6, y = 3\}$

If we get a **weaker formula** as a query, say $(x + y < 10)$, then we can immediately **reuse the solution already found in the cache**, without calling the SMT solver.

If we get a **stronger formula** as a query, say $(x + y < 10) \wedge (x > 5) \wedge (y \geq 0)$, then we can quickly **try the solution in the cache** and see if it works, without calling the solver (in this example, it works).

When constraint solving fails

Despite best efforts, the program may be using constraints in a fragment which the SMT solver does not handle (well).

For example, the SMT solver does not handle **non-linear constraints** well.

When constraint solving fails - example

```
int twice(int v) {  
    return v * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x, y);  
}
```

Here, we changed the twice() function to contain **a non-linear** result.

Let us see what happens when we symbolically execute the program now...

source: Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.

When constraint solving fails - example

```
int twice(int v) {  
    return v * v;  
}  
  
void test(int x, int y) {  
    z = twice(y);  
    if (x == z) {  
        if (x > y + 10)  
            ERROR;  
    }  
}  
  
int main() {  
    x = read();  
    y = read();  
    test(x, y);  
}
```

This is the result if $x = z$:

$\sigma_s : \begin{array}{l} x \mapsto x0, \\ y \mapsto y0 \\ z \mapsto y0*y0 \end{array}$

$pct : x0 = y0*y0$

Now, if we are to invoke the SMT solver with the pct formula, it would be **unable** to compute satisfying assignments, precluding us from knowing whether the path is feasible or not.

source: Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.

- Lecture Notes on "Program Analysis", ETH Zurich, Martin Vechev.
- Lecture Notes on "Techniques for Program Analysis and Verification", Stanford, Clark Barrett.
- Lecture Notes on "Program verification", ETH Zurich, Alexander Summers.
- Lecture Notes on "Computer-Aided Reasoning for Software Engineering", University of Washington, Emina Torlak.